

The background image shows a modern, open-plan office environment. Numerous employees are seated at their desks, each equipped with multiple computer monitors. The office is characterized by its industrial-style ceiling, which features a complex network of white pipes, ducts, and rectangular fluorescent light fixtures. In the foreground, the Pivotal logo is overlaid in white.

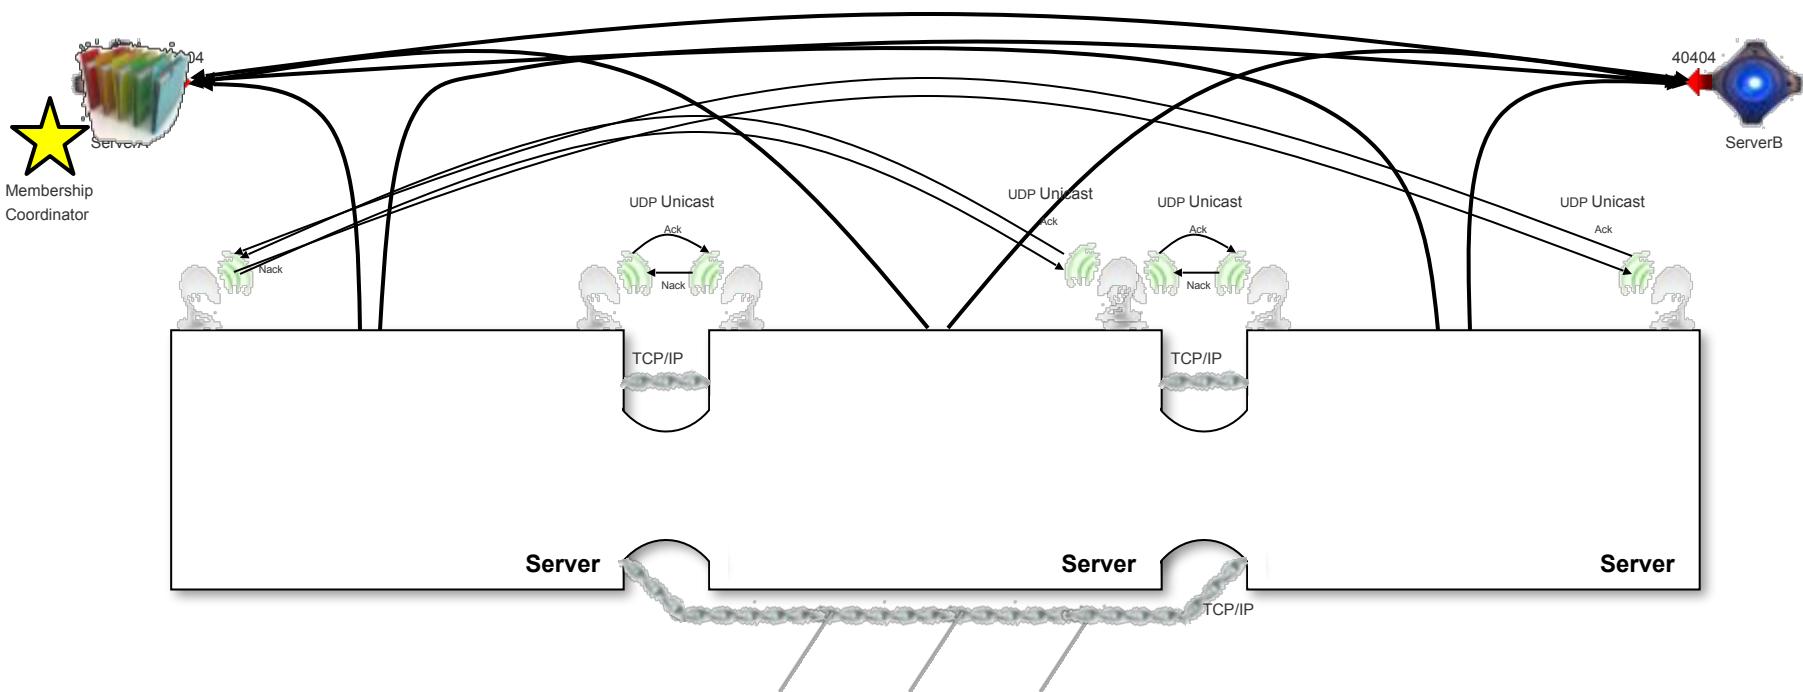
Pivotal

GemFire Technical Background for Operators

Distributed Membership

GemFire Distributed Membership System

```
gemfire.properties  
mcast-port= 0  
locators=ServerA[40404],ServerB[40404]
```



Demonstration

- Start a cluster
- Observe locator and data node logs during normal startup

grep for View in the logs to see normal startup

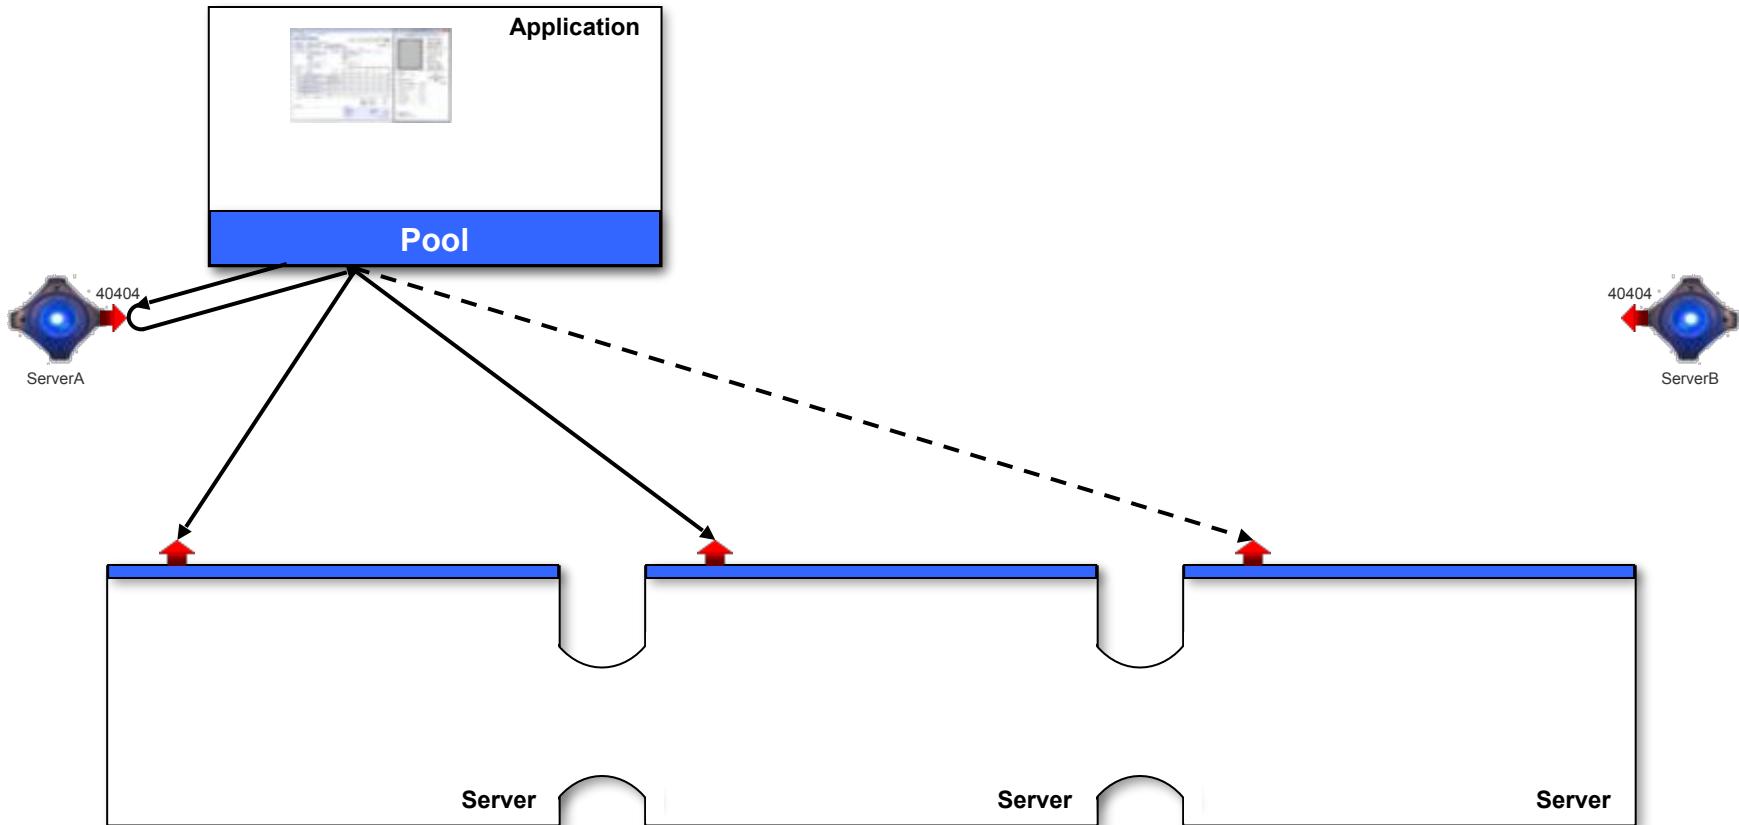
understand how members are identified:

*e.g. 192.168.1.136(locator1:20517:locator)<ec><v0>:10901
host (member-name:PID:member-type)<?><?>:port*

- Kill a member and watch how the distributed membership system reacts

Locators

Clients and Locators



Notes on Locators

- Think of locators as a “cluster information service”. They can answer questions like:
 - What data nodes are in the cluster ?
 - Which data node holds a certain bucket of data ?
 - What port is that data node listening on ?
- Servers also have a locators setting because when they come up they register information with the locators.
- Locators allow servers to come and go without reconfiguring the clients.
- Data does not flow through them. You can even live without them (as long as nothing changes).
- Everything in the cluster (locators and data nodes) should have the same “locators” setting
- “in the same cluster” is the same as “has the same locators setting”

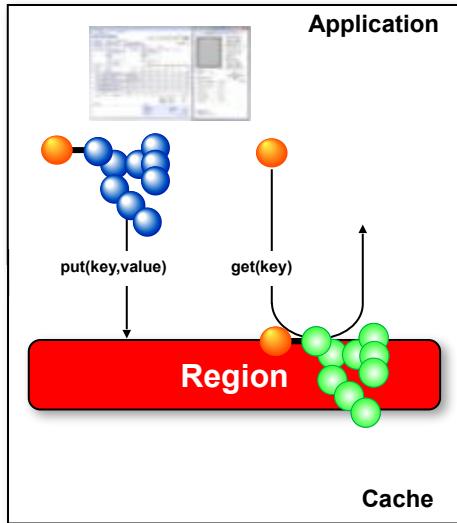
Demonstration

We'll show that you can use a cluster even when all locators are temporarily stopped.

How GemFire Manages Data

Data Management

A region is a container for similar objects ~ table



```
// Insert or update a customer record into the cache  
customerRegion .put(customerIdX, customerObjX);
```

```
// Alternatively, insert or update a customer record into the cache with a context  
customerRegion .put(customerIdX, customerObjX, context);
```

```
// Retrieve a customer record from the cache  
Customer customerY = customerRegion .get(customerIdY);
```

```
// Retrieve a customer record from the cache with a context  
Customer customerY = customerRegion .get(customerIdY, context);
```

Regions also implement [Map<K,V>](#), [ConcurrentMap<K,V>](#)

Region Types

Regions come in 2 main “flavors”

REPLICATE

- All data is in every data node
- When a change occurs it is replicated to all data nodes synchronously
- Write performance gets worse as the cluster grows

PARTITION

- Data is “striped” across data nodes
- Backups are striped too but a bucket and its backup will not live on the same data node
- When a change occurs it is replicated to one (or two) backup(s) on another node
- Scales linearly

Partition Region Concept: Buckets and Routing

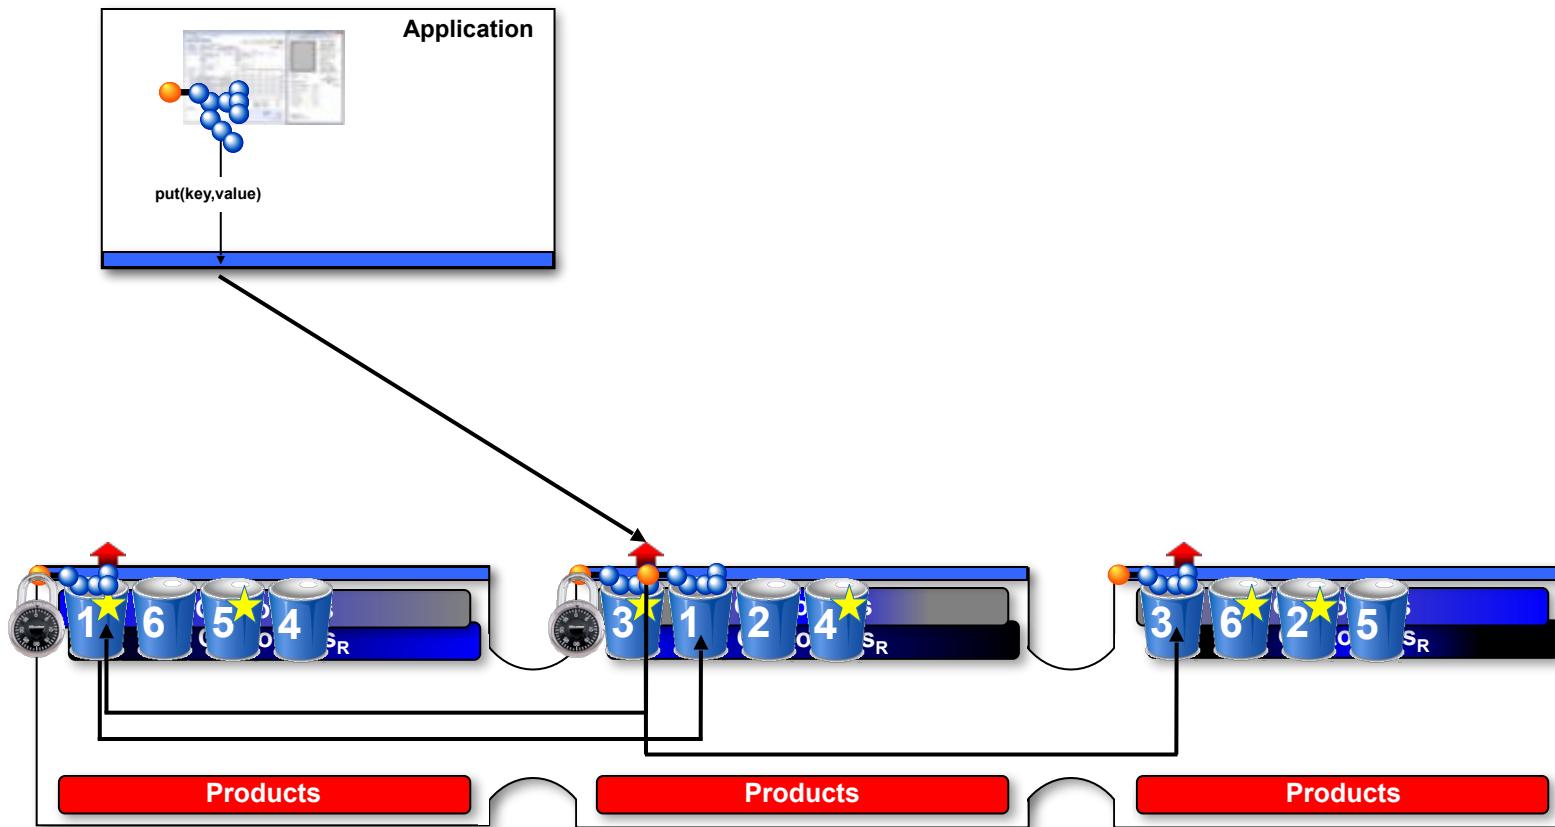


Partition Redundancy

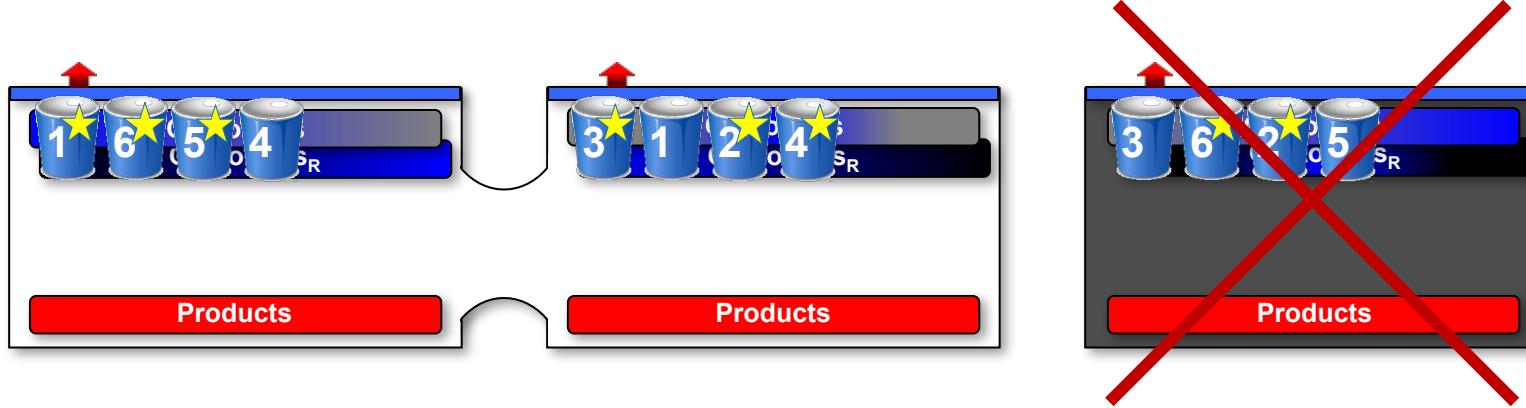
```
// gfsh  
create region --name=Customers --type=PARTITION_REDUNDANT --redundant-copies=1
```



Partition Redundancy

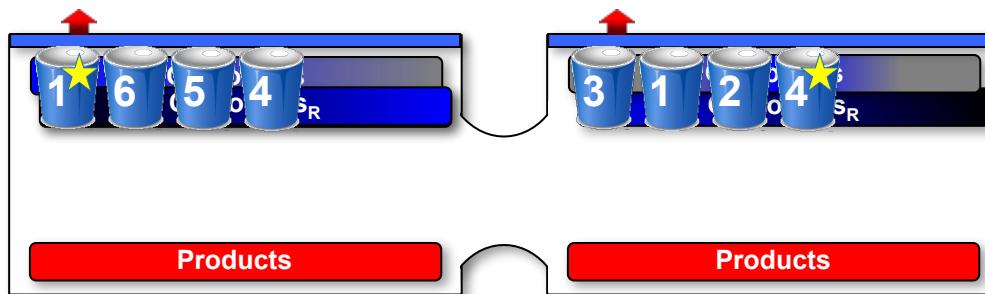


Partition Redundancy



Partition Regions - Recovery Delay

```
//gfs  
create region --name=Customers --type=PARTITION_REDUNDANT --redundant-copies=2 --recovery-delay=10000
```



Partition Regions - Startup Recovery Delay

```
//gfs  
create region --name=Customers --type=PARTITION_REDUNDANT --redundant-copies=1 ---recovery-delay=10000 --startup-recover-delay=10000
```

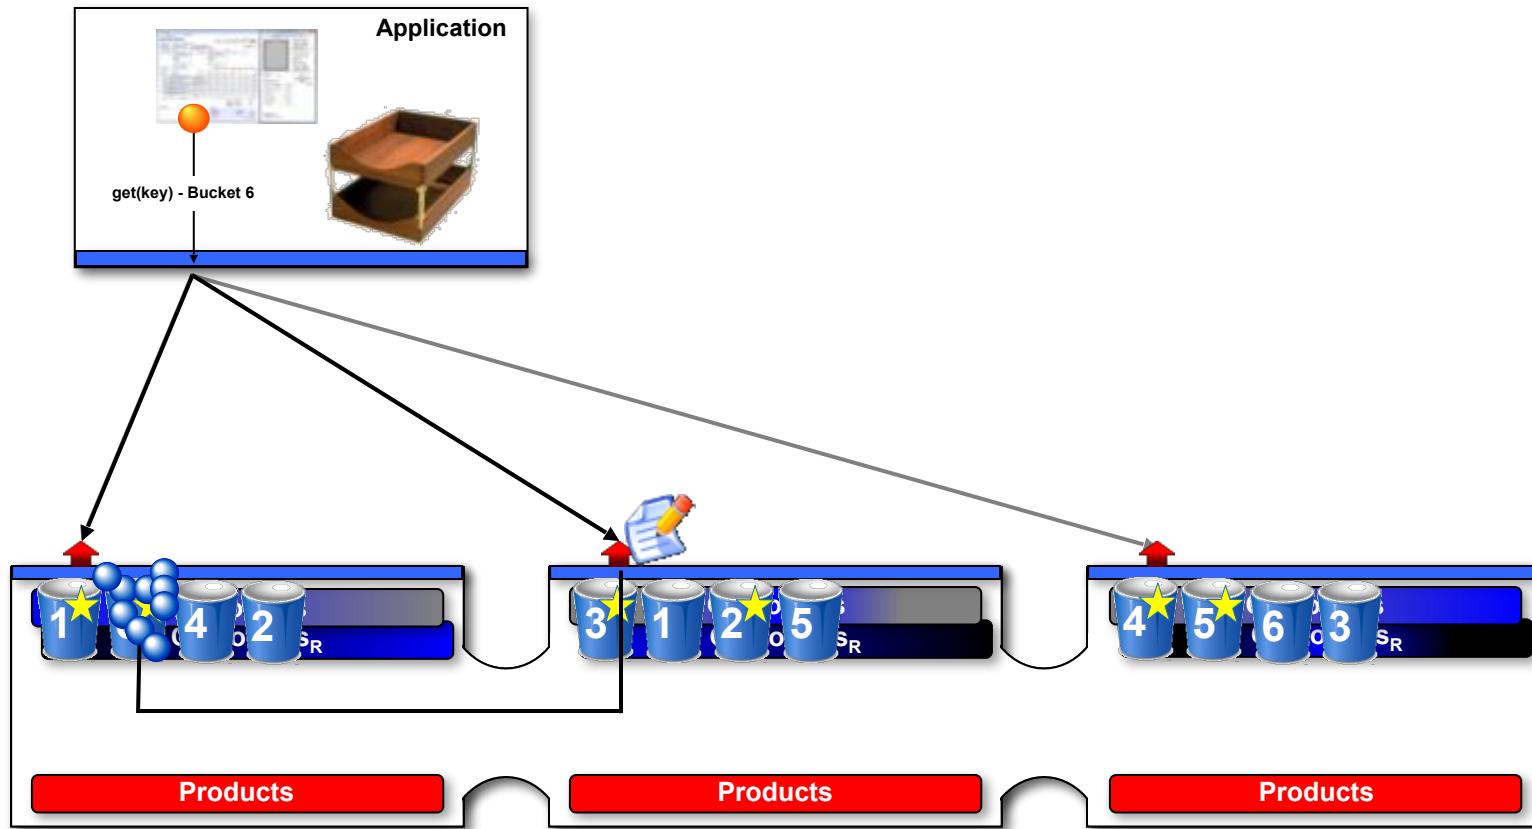


Colocate Regions

```
// gfsh  
create region --name=Orders --type=PARTITION_REDUNDANT --redundant-copies=1 --colocated-with=Customers
```



Single-Hop Optimization



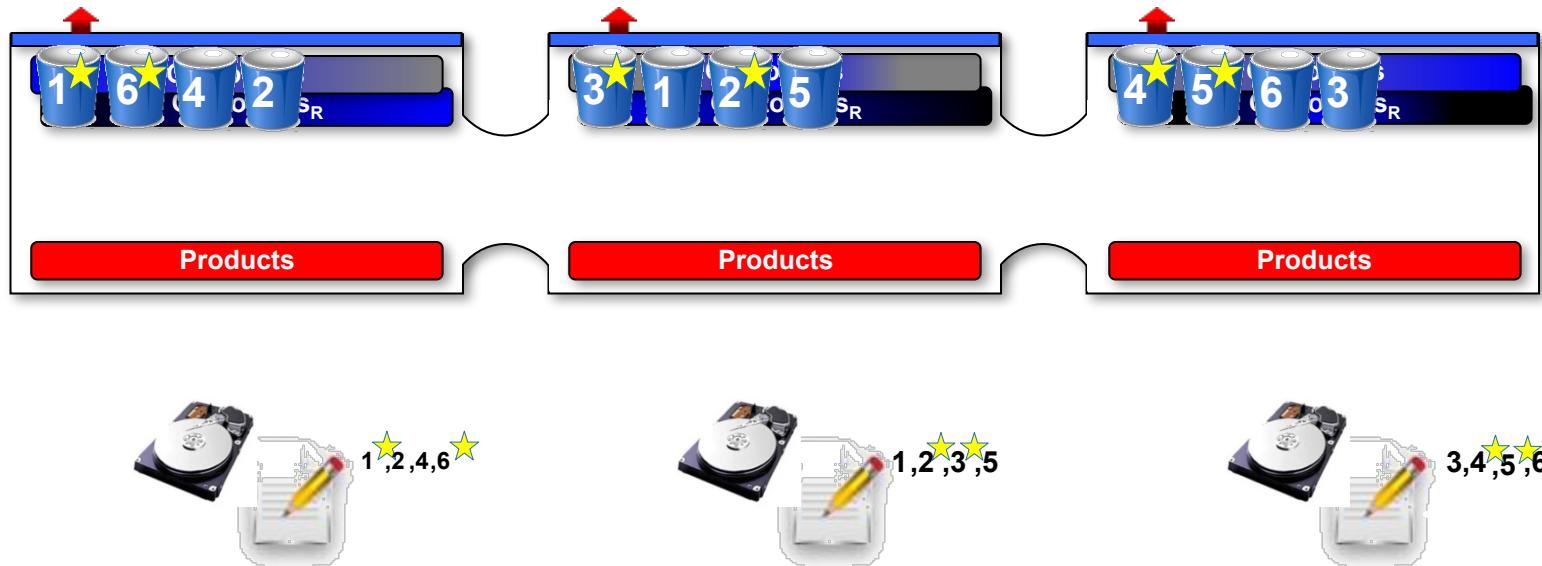
Demonstration

We'll watch partitioned regions and redundancy recovery in action.

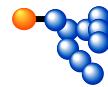
Parallel, Shared Nothing Persistence



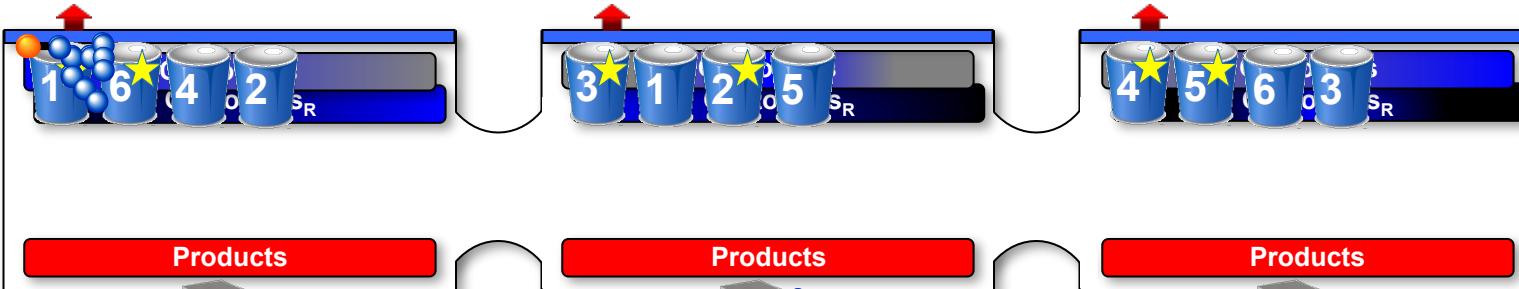
Parallel, Shared Nothing Persistence



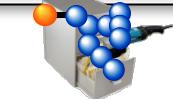
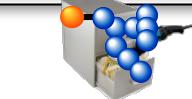
Parallel, Shared Nothing Persistence



Locator
(Membership Coordinator)



OS File Buffer



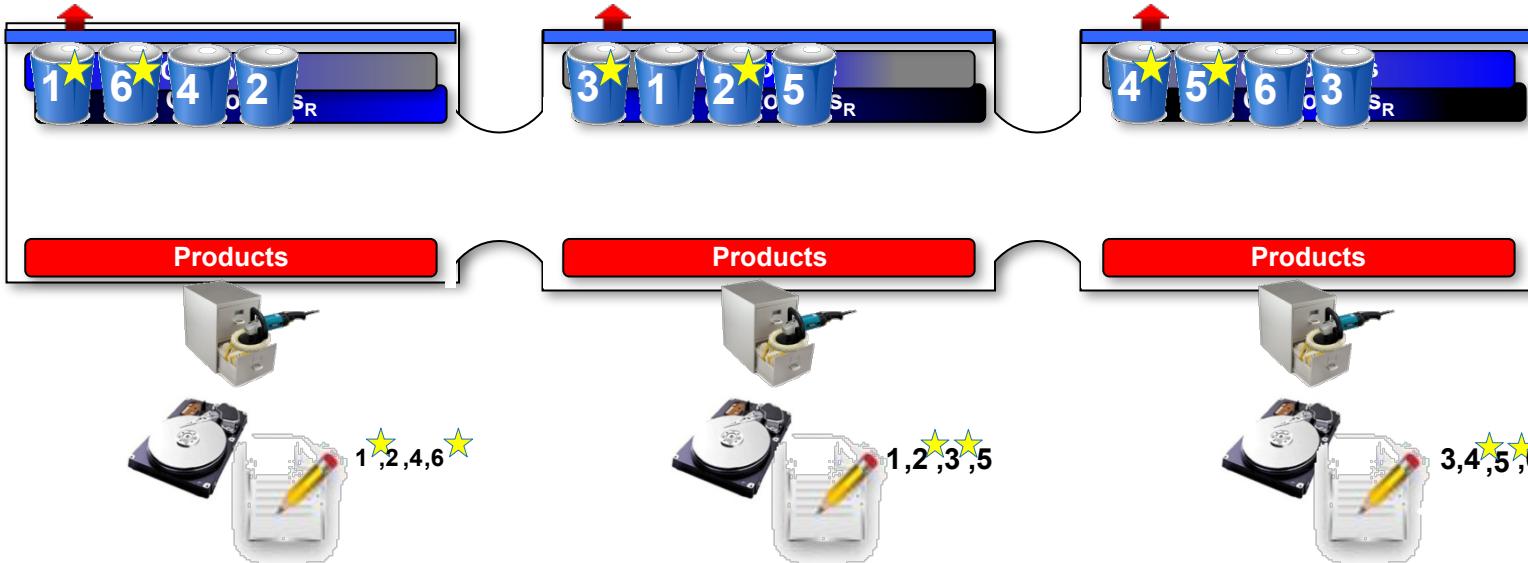
1,2,4,6

1,2,3,5

3,4,5,6

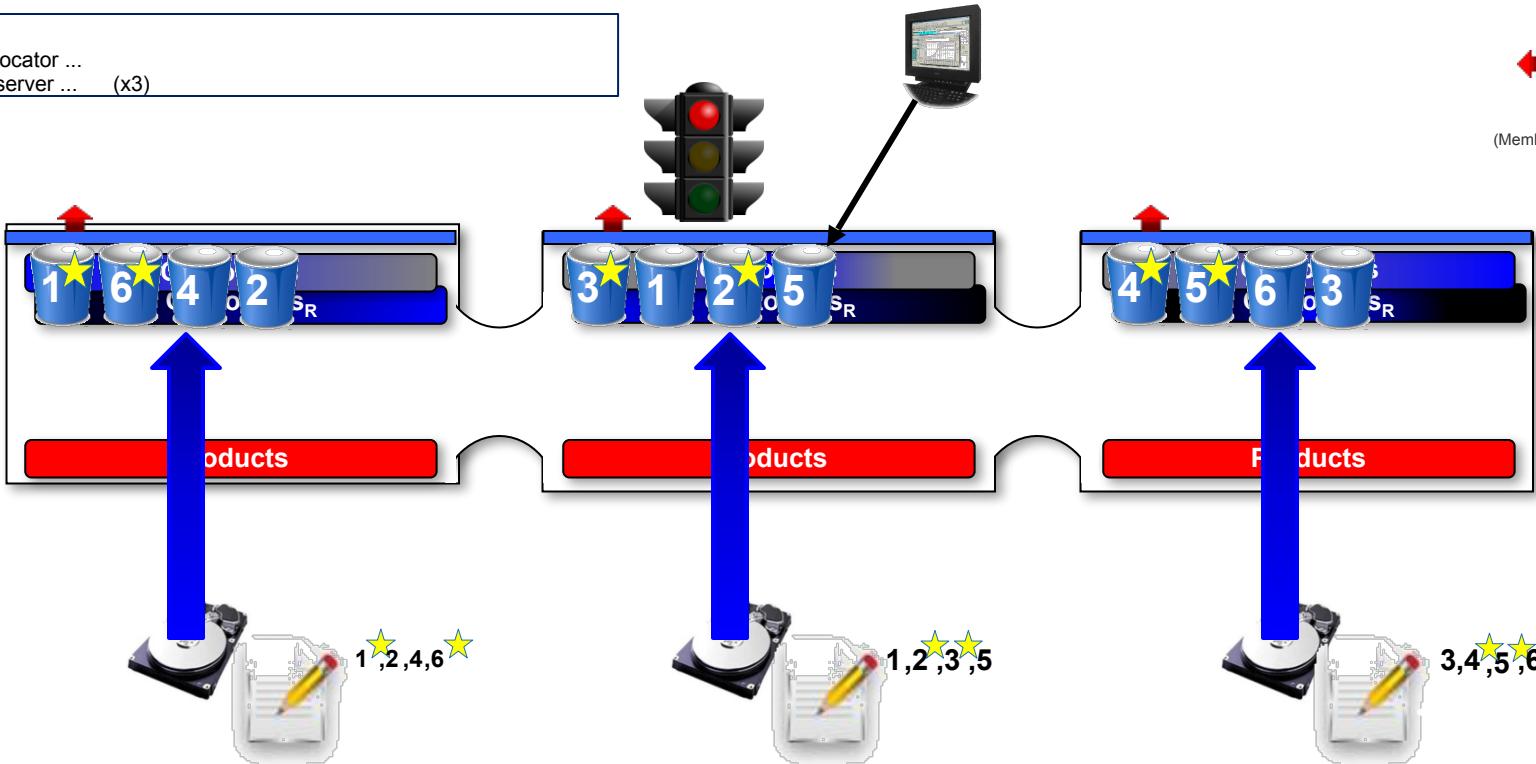
Shutdown

```
> gemfire shut-down-all
```



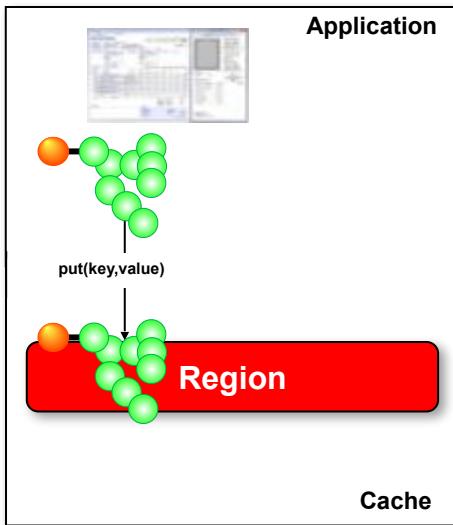
Parallel Recovery

```
> gfsh start locator ...  
> gfsh start server ... (x3)
```



Disk Store Details

Disk Stores



```
//gfsh  
create disk-store --name=persistentStore --dir=/data1/customers#1000 --dir=/data2/customers  
--dir=/data3/customers --auto-compact=true --allow-force-compaction=true --max-oplog-size=100
```

-rw-rw-r-- 1 gemfire users 1924 Mar 22 13:57 BACKUPpersistentStore_if
-rw-rw-r-- 1 gemfire users 0 Mar 22 13:56 DRLK_IFpersistentStore_if
-rw-rw-r-- 1 gemfire users 943718 Mar 22 13:57 BACKUPpersistentStore_1cif
-rw-rw-r-- 1 gemfire users 104857 Mar 22 13:57 BACKUPpersistentStore_1dn

90% - pre-allocated - 10% max-oplog-size

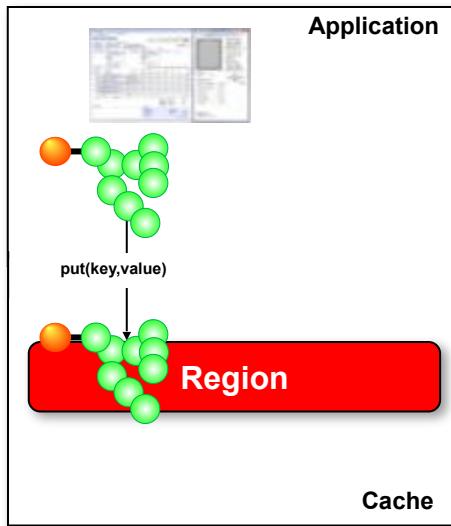
Runtime Lock File

Create, Update, Invalidate Operations

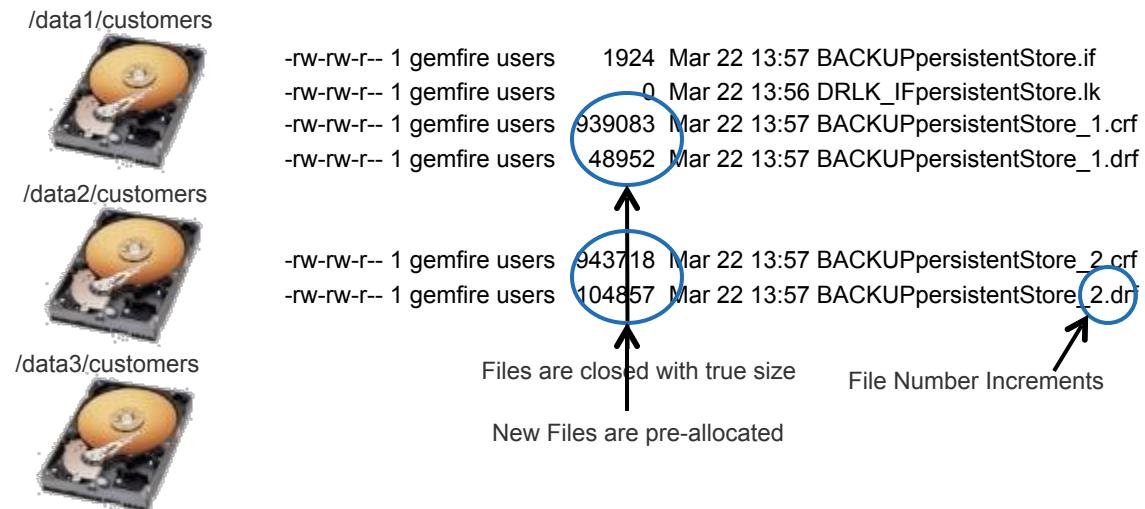
Delete Operations

Disk Stores - Rollover

```
//gfsh create disk-store --name=persistentStore --dir=/data1/customers#1000 --dir=/data2/customers  
--dir=/data3/customers --auto-compact=true --allow-force-compaction=true --max-oplog-size=100
```

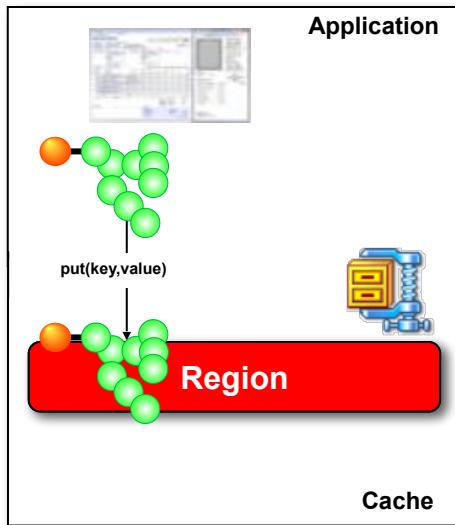


When oplogs reach max-oplog-size



Disk Stores - Compaction

```
//gfsh  
create disk-store --name=persistentStore --dir=/data1/customers#1000 --dir=/data2/customers  
--dir=/data3/customers --auto-compact=true --allow-force-compaction=true --max-oplog-size=100
```



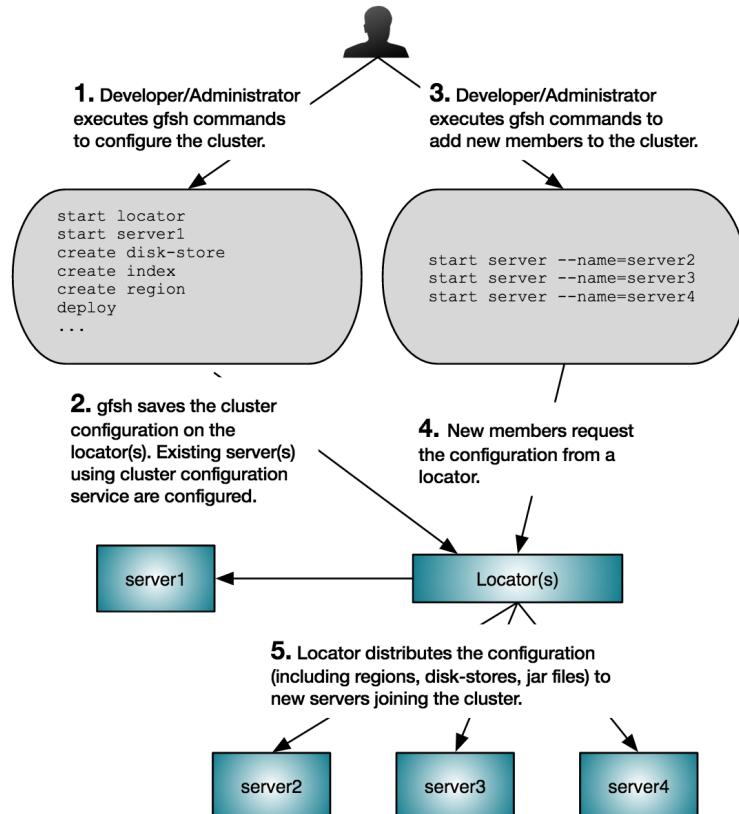
When auto-compact is true or by invoking
`compact disk-store --name=persistentStore`



```
-rw-rw-r-- 1 gemfire users      1924 Mar 22 13:57 BACKUPpersistentStore.if  
-rw-rw-r-- 1 gemfire users       0 Mar 22 13:56 DRLK_IFpersistentStore.lk  
-rw-rw-r-- 1 gemfire users  939083 Mar 22 13:57 BACKUPpersistentStore_1.crf  
-rw-rw-r-- 1 gemfire users   48952 Mar 22 13:57 BACKUPpersistentStore_1.drf  
  
-rw-rw-r-- 1 gemfire users  912428 Mar 22 13:57 BACKUPpersistentStore_2.crf  
-rw-rw-r-- 1 gemfire users  104687 Mar 22 13:57 BACKUPpersistentStore_2.drf  
  
-rw-rw-r-- 1 gemfire users  943718 Mar 22 13:57 BACKUPpersistentStore_3.crf  
-rw-rw-r-- 1 gemfire users  104857 Mar 22 13:57 BACKUPpersistentStore_3.drf
```

Cluster Configuration Service

The Cluster Configuration Service - How It Works



The Cluster Configuration Service - Caveats and Best Practices

- **Do Not Use With cache.xml**

This will appear to work but can lead to problems that prevent nodes from coming up after restart.

- **Data is stored in the locator working directory - do not delete it !**
- **PDX configuration does not apply to running members.**

This leads to a chicken and egg problem with new clusters that will use pdx. To overcome this, when a new cluster is first created, follow this procedure.

1. start the cluster
2. create the pdx disk store: **gfsh> create disk-store --name=pdx-disk-store --dir=/data/pdx**
3. configure pdx: **configure pdx --disk-store=pdx-disk-store --read-serialized=true**
4. stop all the data nodes (locators can stay up): **gfsh> shutdown**
5. start the cluster

Demonstration: Cluster Configuration Service

We will see ...

- that configuration made with gfsh is persisted across restarts
- where the configuration is stored
- how to tell what configuration is being used

GemFire and the File System

GemFire Files

Product Directory

- Can be anywhere
- Can be read only

Working Directory

- Must be writable by “run as” user
- Required by both locators and servers
- Set with the “--dir” option of the “gfsh start server/locator” command

Logs and Stats

- gemfire automatically rolls logs and stats
- use properties to set location and size limits
- log/stats location must be writable, should be local

Disk Stores (where the data lives)

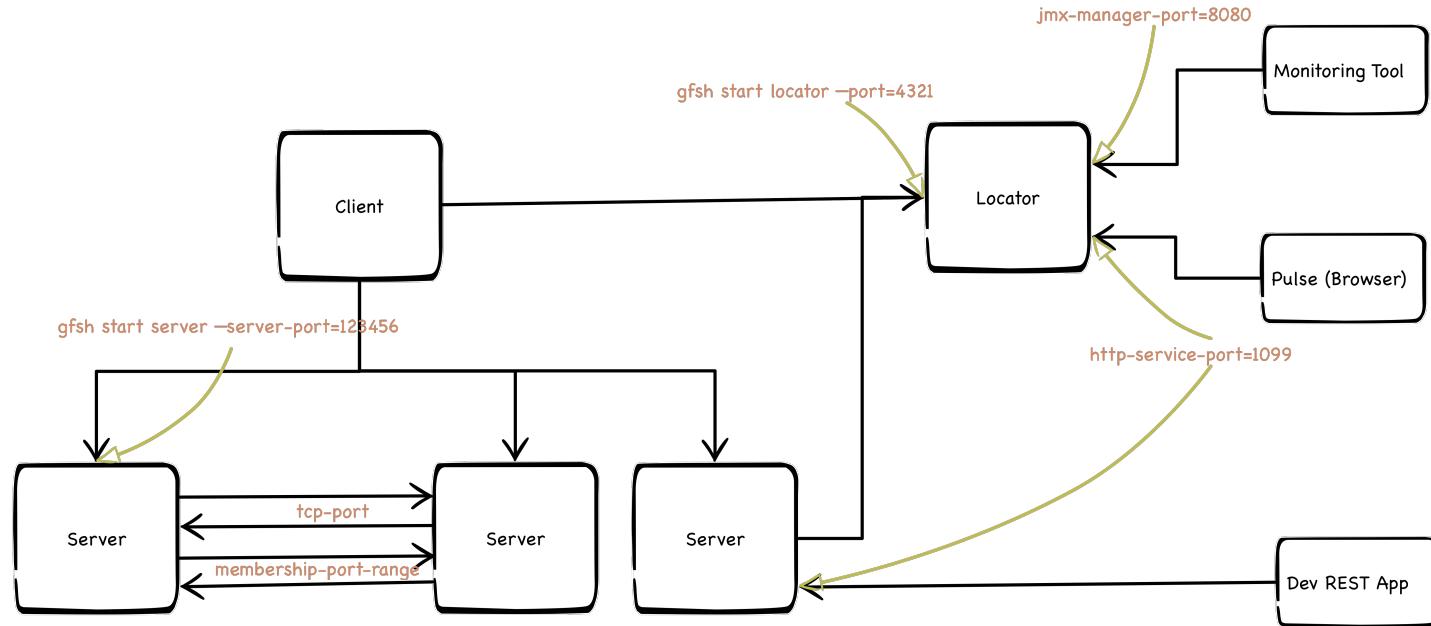
- Must be writeable by “run as” user

GemFire Logs and Archives

Sample Log Configuration

```
log-level=config
statistic-sampling-enabled=true
statistic-archive-file=server.gfs
archive-file-size-limit=10
archive-disk-space-limit=100
log-file=server.log
log-file-size-limit=10
log-disk-space-limit=100
```

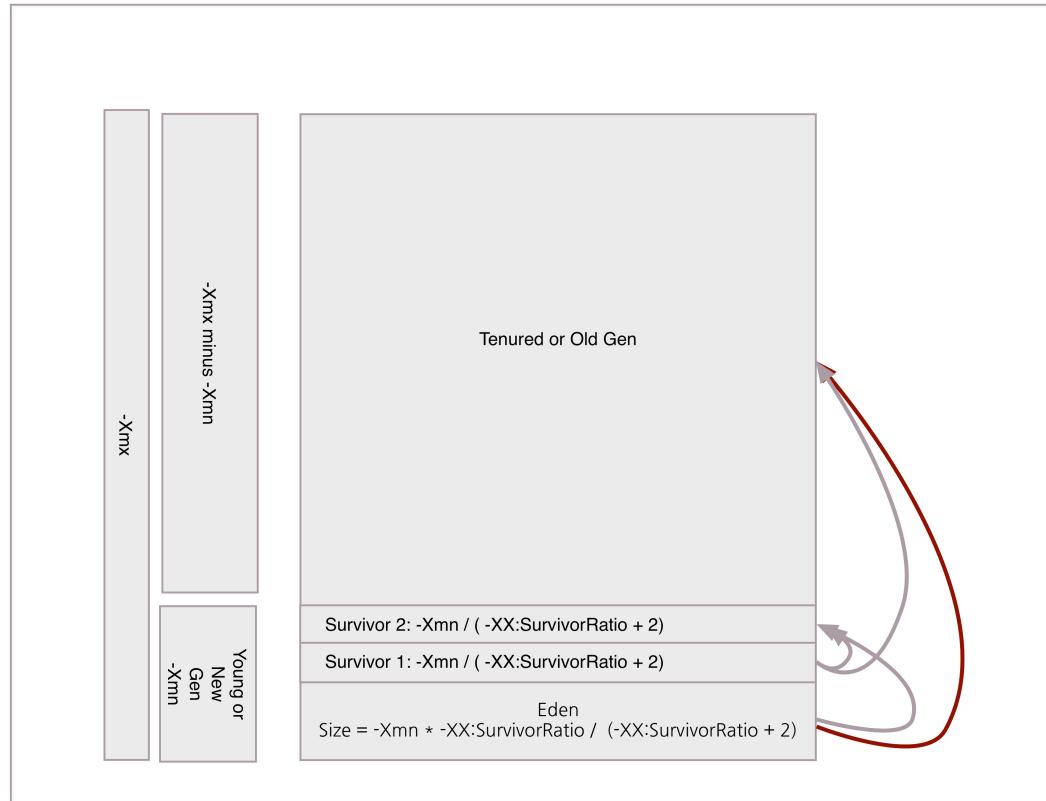
GemFire and the Network



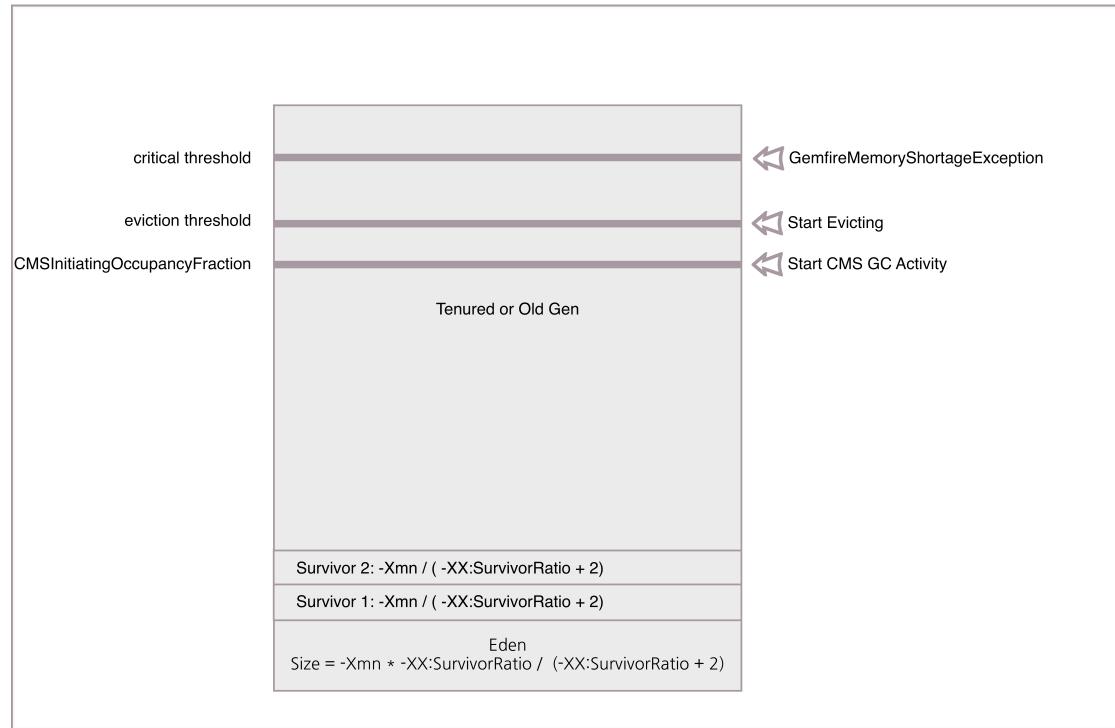
Also Bind Addresses ...
bind-address (peer to peer)
server-bind-address (client-server)
http-service-bind-address
jmx-manager-bind-address

GemFire and Memory Management

GemFire and Memory Management



GemFire and Memory Management



Interaction Between the Resource Manager and Garbage Collection

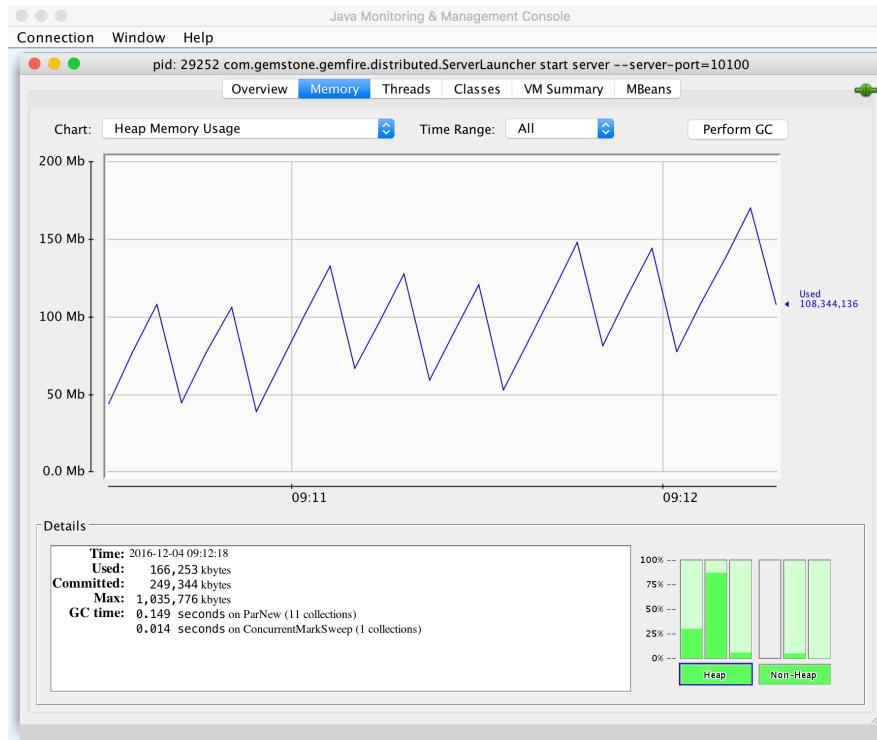
GemFire and the JVM: JVM Option Recommendations

- Maximum and minimum memory should be the same, use parallel new gc collector and cms old collector, set the young gen to roughly 1/8th of total heap but not less than 1G and not more than 8G
- Set the CMSInitiatingOccupancyFraction to slightly less than the eviction threshold.
- Example of a 32G JVM with a 80% eviction threshold

```
--J=-Xmx32g --J=-Xms32g --J=-Xmn4g --J=-XX:+UseConcMarkSweepGC --  
J=-XX:+UseParNewGC --J=-XX:CMSInitiatingOccupancyFraction=75 --J=-  
XX:+UnlockDiagnosticVMOptions --J=-XX:ParGCCardsPerStrideChunk=4096
```

Demonstration: Cluster Configuration Service

We will watch the 2 garbage collectors in action.



Essential GemFire Configuration

Setting Startup Properties

There are *many* startup properties:

http://gemfire.docs.pivotal.io/95/geode/reference/topics/gemfire_properties.html

Properties can be set multiple ways. In order of precedence:

As a java system property on the gfsh start command line

gfsh start server —J=-Dgemfire.locators=myhost[1234] ...

With a supported gfsh argument (only available for some properties)

gfsh start server —locators=myhost[1234]

In a properties file (defaults to “gemfire.properties”)

```
# gemfire.properties  
locators=myhost[1234]
```

Essential GemFire Configuration

Key configuration files

“gemfire.properties”

- Set ports, file locations and many other things
- Settings are not application specific

“gfsecurity.properties”

- Separate file where security and SSL related properties can be placed
- This file can have more restrictive permissions (of course gemfire still needs to read).
- Allows security to be added to an existing server or client. *Note: you can use it with gfsh*

“cache.xml”

- Defines regions, disk stores
- Most settings are application specific
- *This has been superseded by gfsh and the region configuration service. You usually don't need it.*

How GemFire Finds its Property Files

see: http://gemfire.docs.pivotal.io/95/geode/configuring/running/default_file_specs.html

Default File Name	Search Path	Java System Property to Override Search Path
gemfire.properties	1. Current Directory 2. User Home Directory 3. Class Path	<code>-J=-DgemfirePropertyFile=X</code>
gfsecurity.properties	1. Current Directory 2. User Home Directory 3. Class Path	NA
cache.xml	1. Current Directory 2. Class Path	<code>cache-xml-file=x</code> in gemfire.properties OR <code>--cache-xml-file=x</code> on gfsh start server OR <code>--J=-Dgemfire.cache-xml-file=x</code> on gfsh start server

Essential GemFire Configuration

Key References

official documentation:

<http://gemfire.docs.pivotal.io>

gfsh:

http://gemfire.docs.pivotal.io/docs-gemfire/latest/tools_modules/gfsh/chapter_overview.html

gemfire.properties and cache.xml:

http://gemfire.docs.pivotal.io/docs-gemfire/latest/reference/book_intro.html

Pivotal[®]

Transforming How The World Builds Software